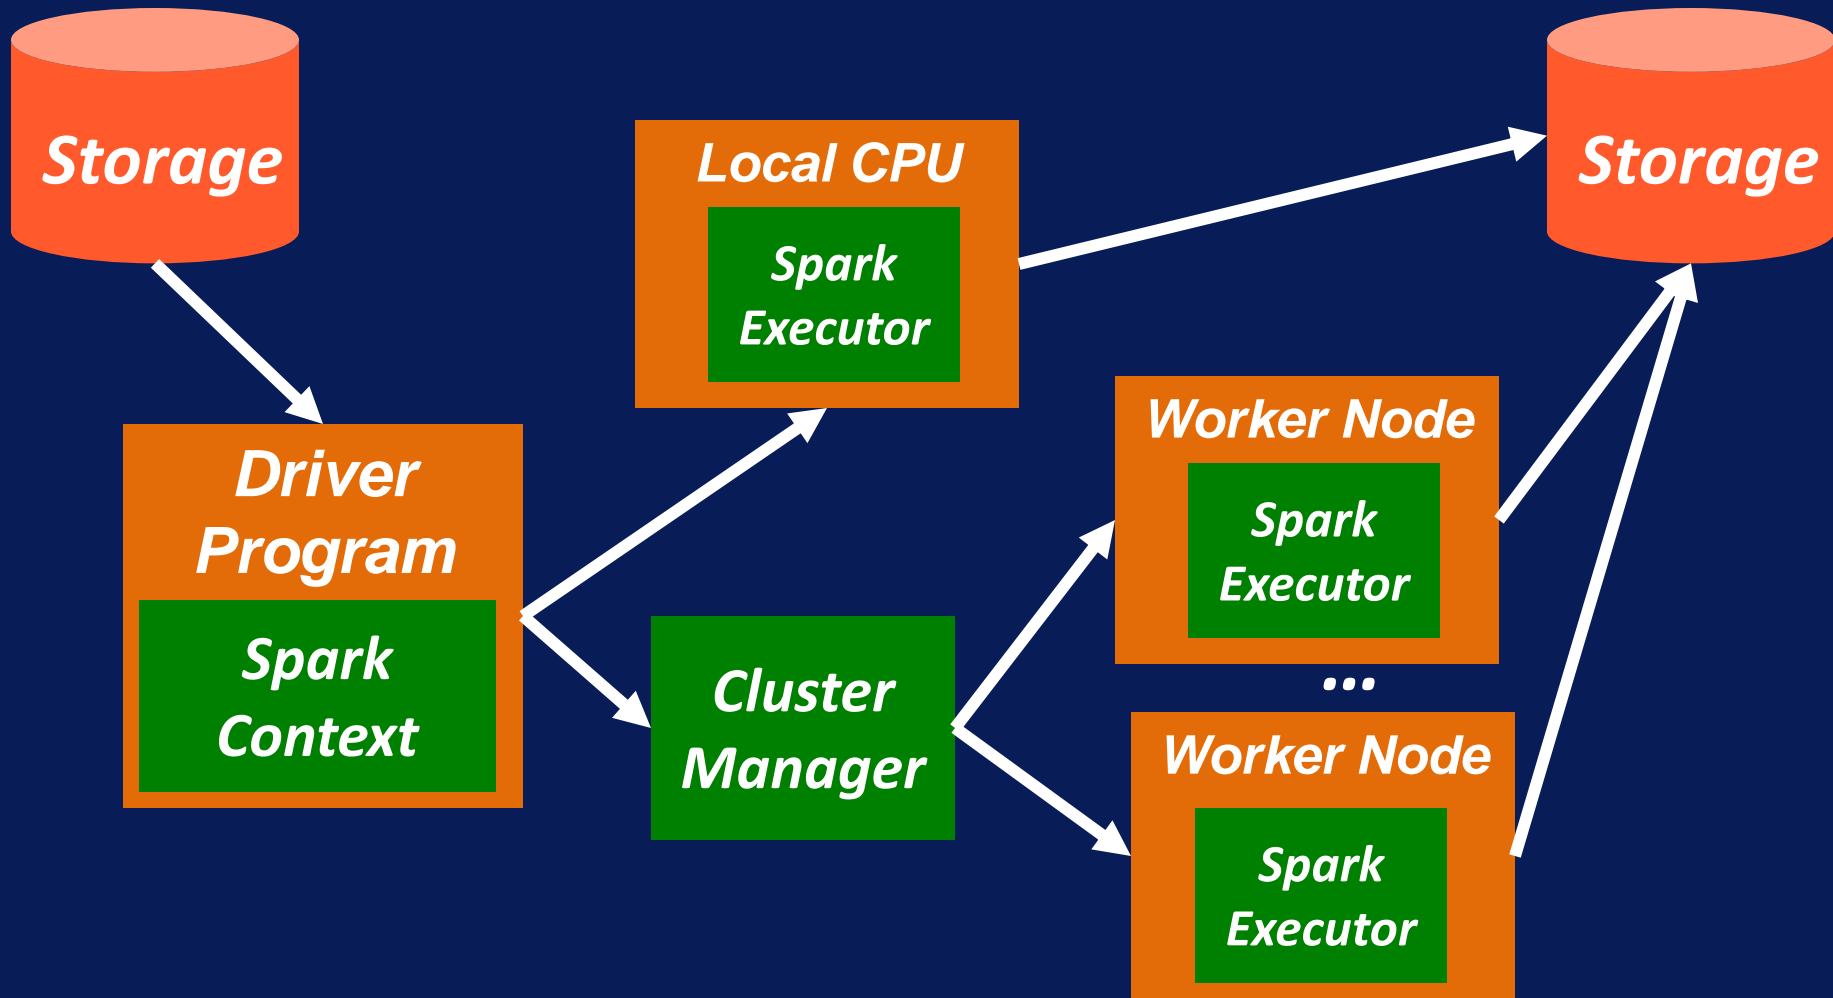


Spark Core: Programming In Spark



After this video you will be able to..

- Use two methods to create RDDs in Spark
- Explain what immutable means
- Interpret a Spark program as a pipeline of transformations and actions
- List the steps to create a Spark program



Creating RDDs

*Driver
Program*

```
In [1]: lines = sc.textFile("hdfs://user/cloudera/words.txt")
```

```
lines = sc.parallelize(["big", "data"])
```

```
numbers = sc.parallelize(range(10), 3)
```

Parallelize
range output
into 3 partitions

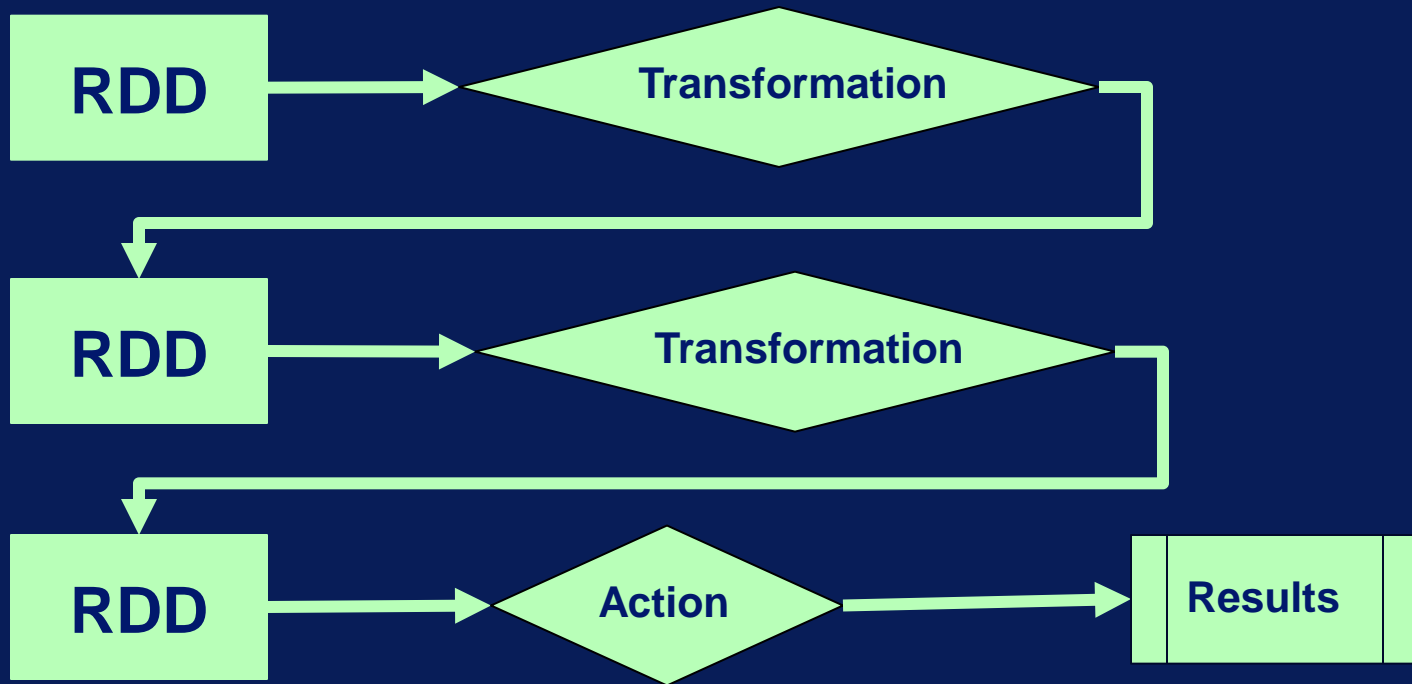
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]


```
numbers.collect()
```

[0, 1, 2], [3, 4, 5], [6, 7, 8, 9]

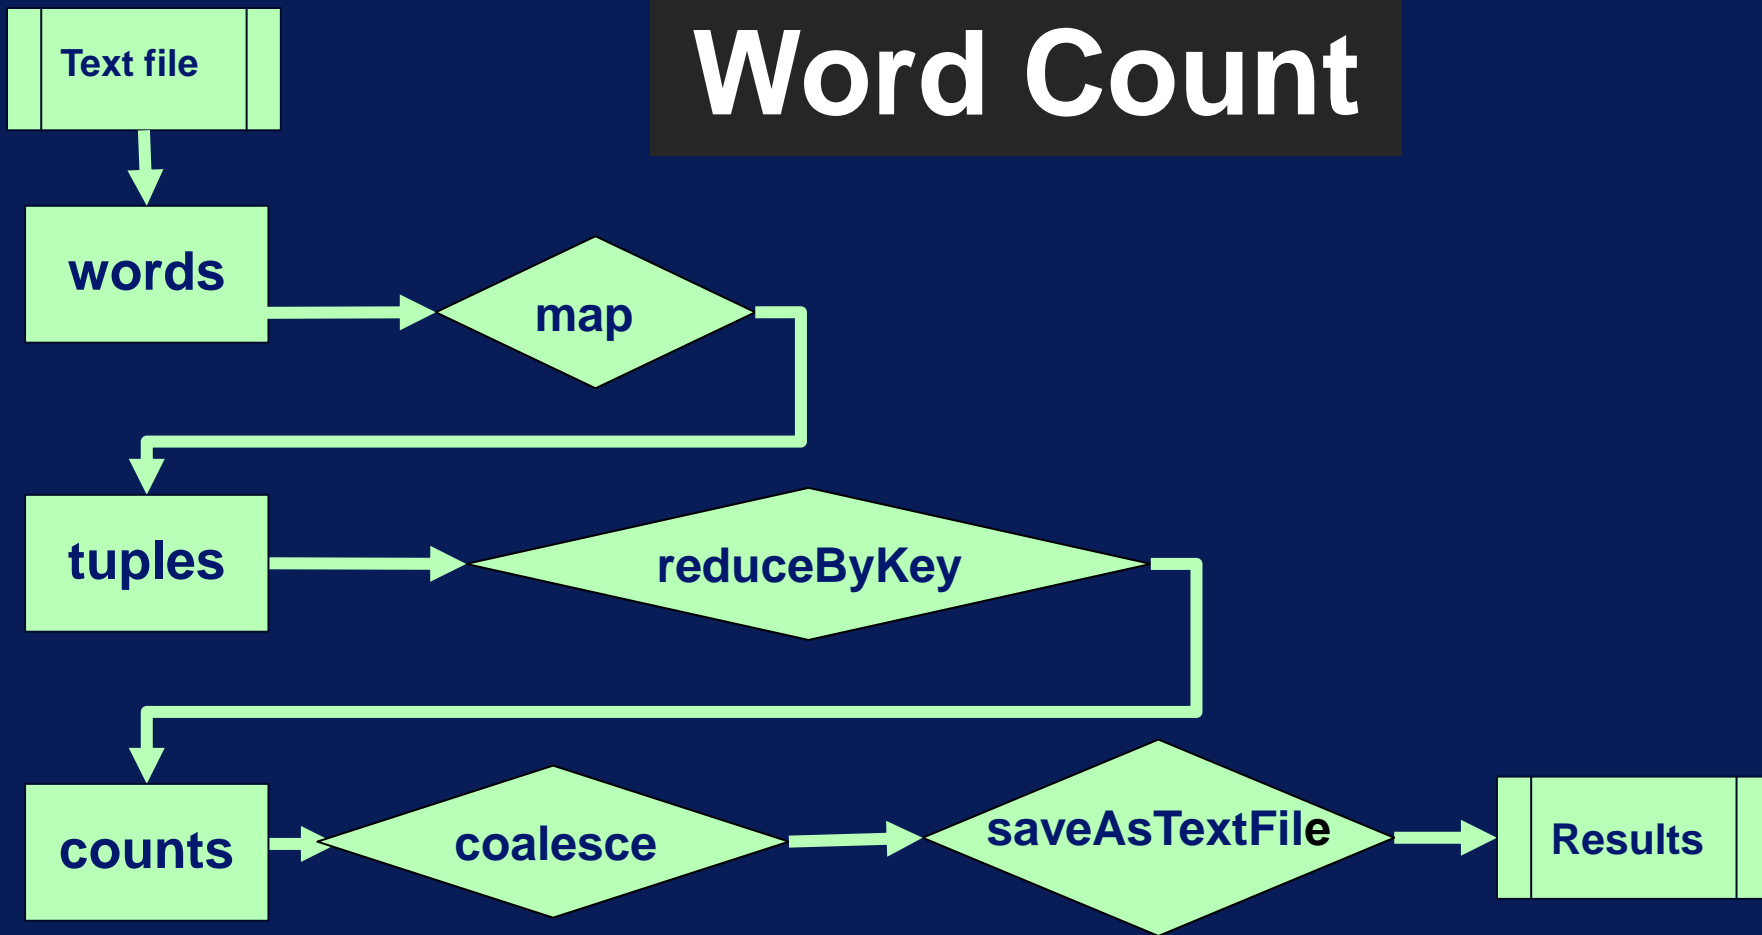
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]

Processing RDDs





Word Count



Programming in Spark

Create RDDs



Apply transformations



Perform actions

Spark Core: Transformations

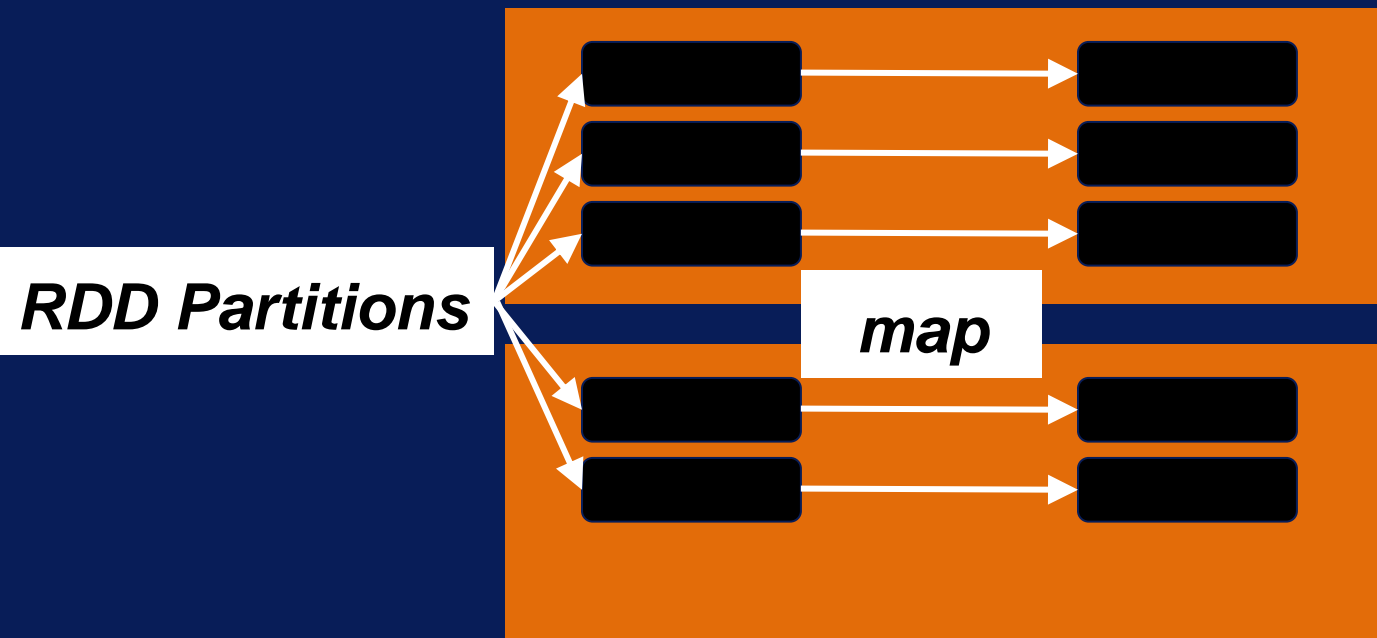


After this video you will be able to..

- Explain the difference between a narrow transformation and wide transformation
- Describe map, flatmap, filter and coalesce as narrow transformations
- List two wide transformations

map

map : apply function to each element of RDD



map

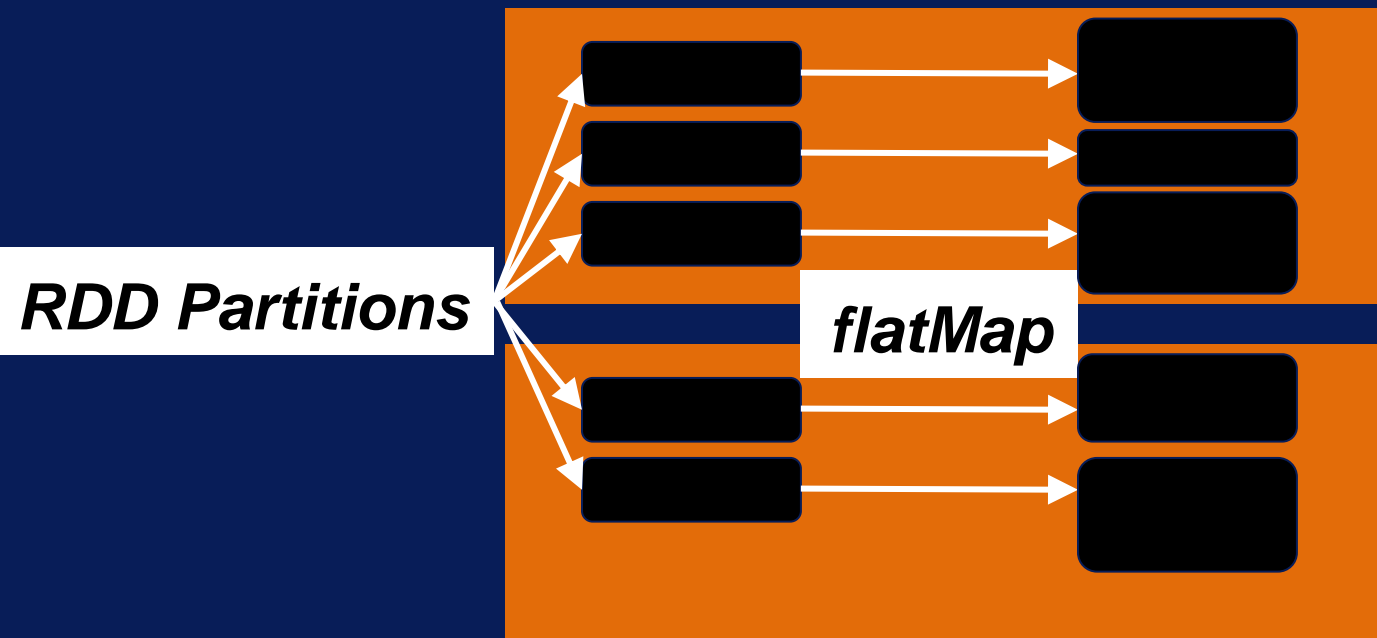
map: apply function to
each element of RDD



```
def lower(line):  
    return line.lower()  
  
lower_text_RDD = text_RDD.map(lower)
```

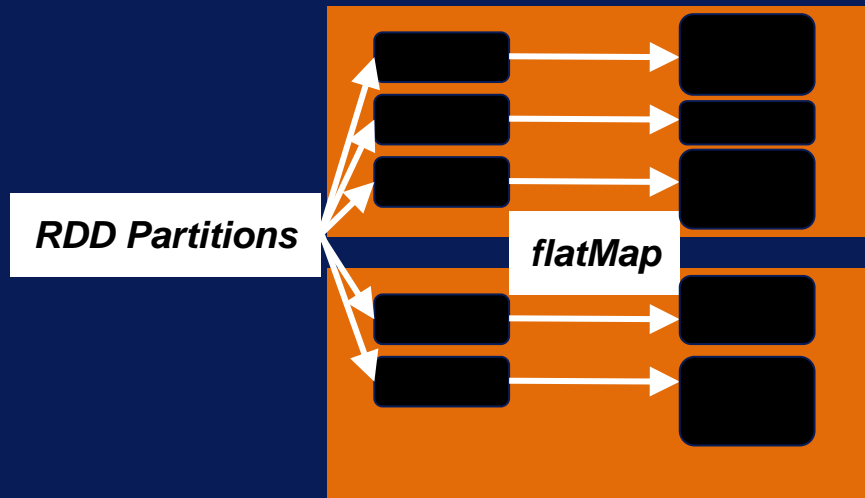
flatMap

flatMap : map then flatten output



flatMap

flatMap: map then
flatten output

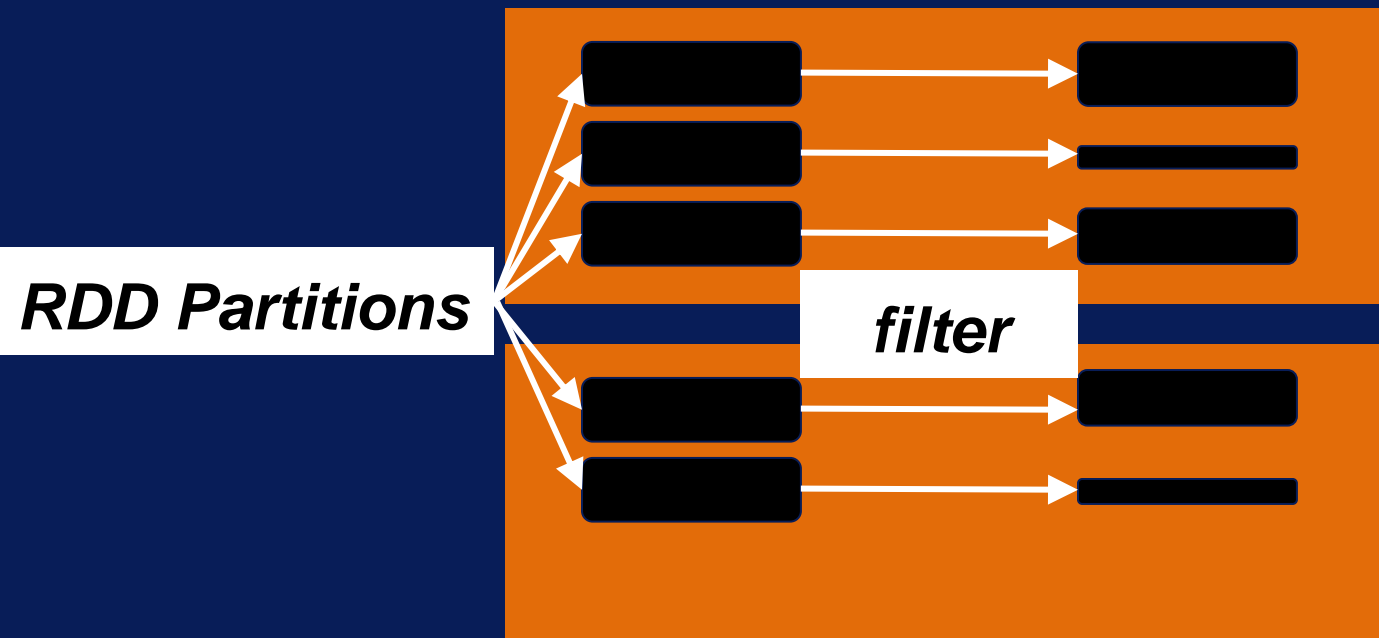


```
def split_words(line):  
    return line.split()
```

```
words_RDD = text_RDD.flatMap(split_words)  
words_RDD.collect()
```

filter

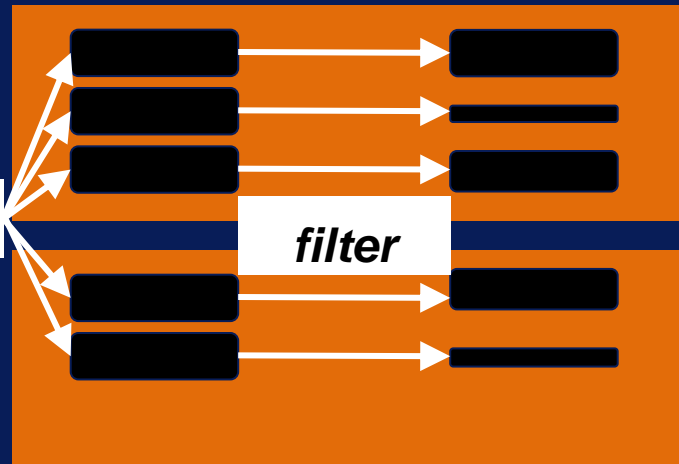
filter : keep only elements where function is true



filter

filter: keep only elements
where function is true


RDD Partitions



```
def starts_with_a(word):  
    return word.lower().startswith("a")  
  
words_RDD.filter(starts_with_a).collect()
```

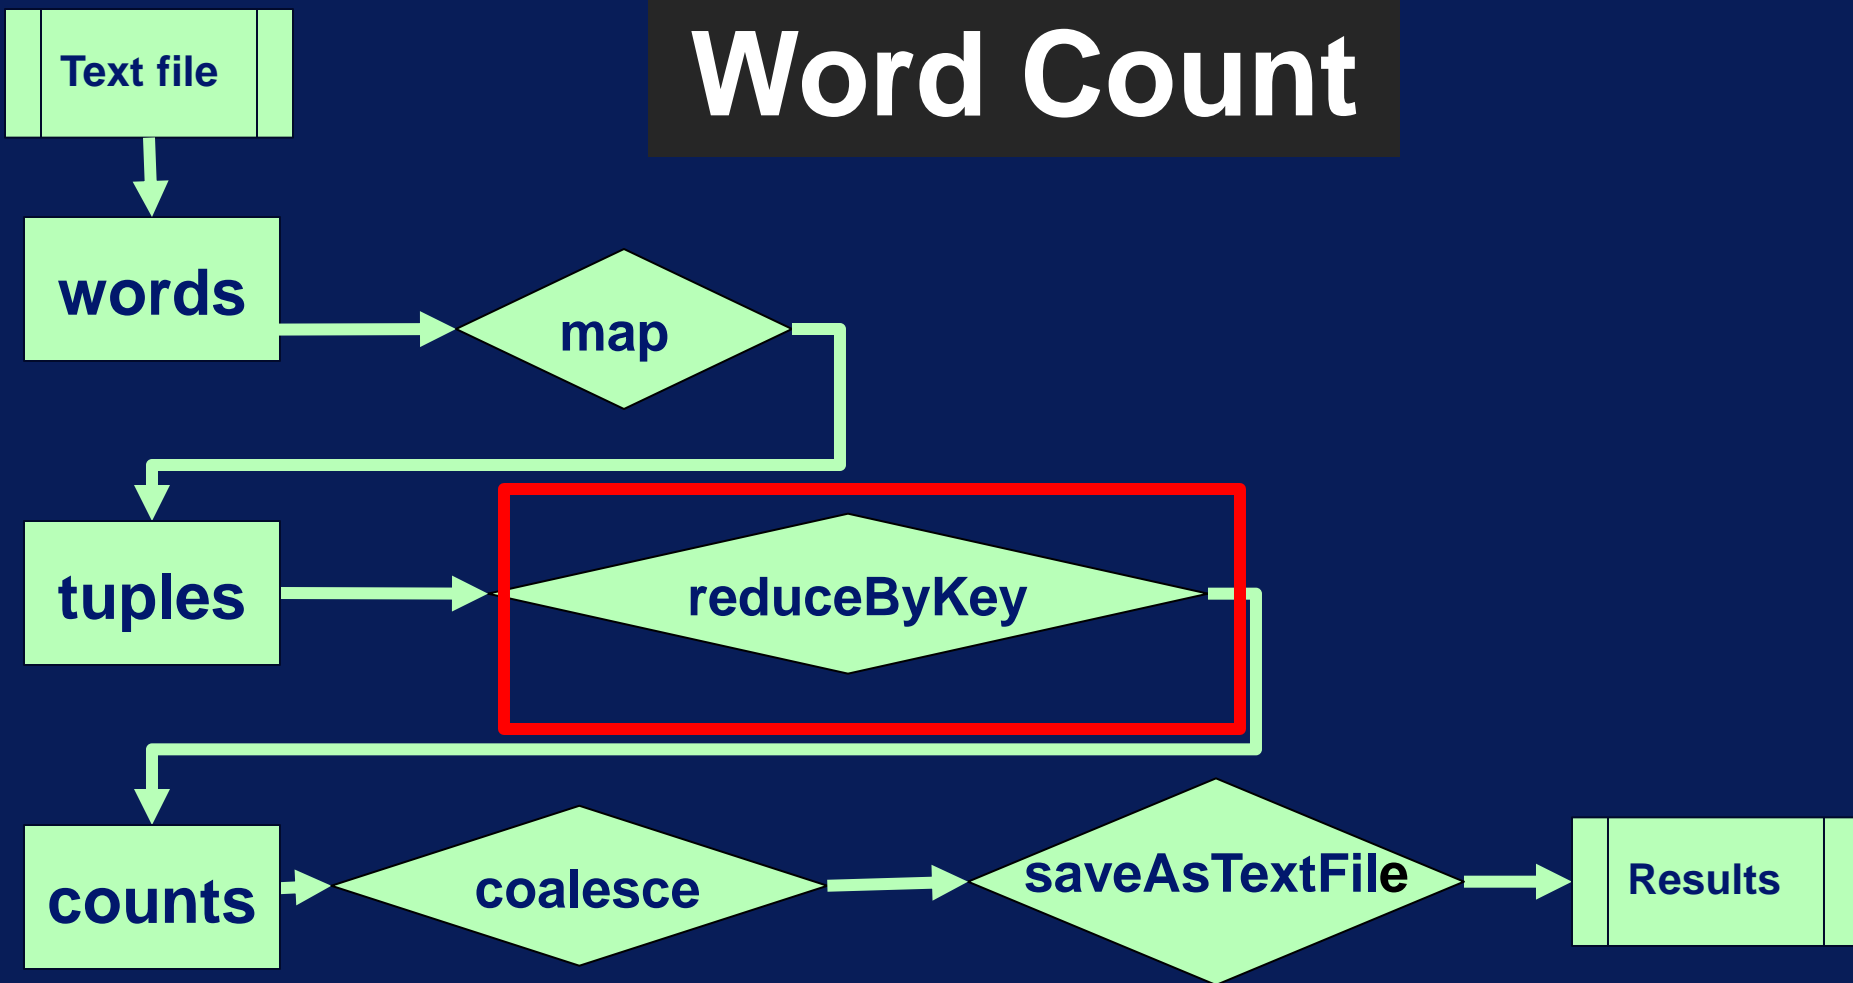

coalesce

coalesce : reduce the number of partitions

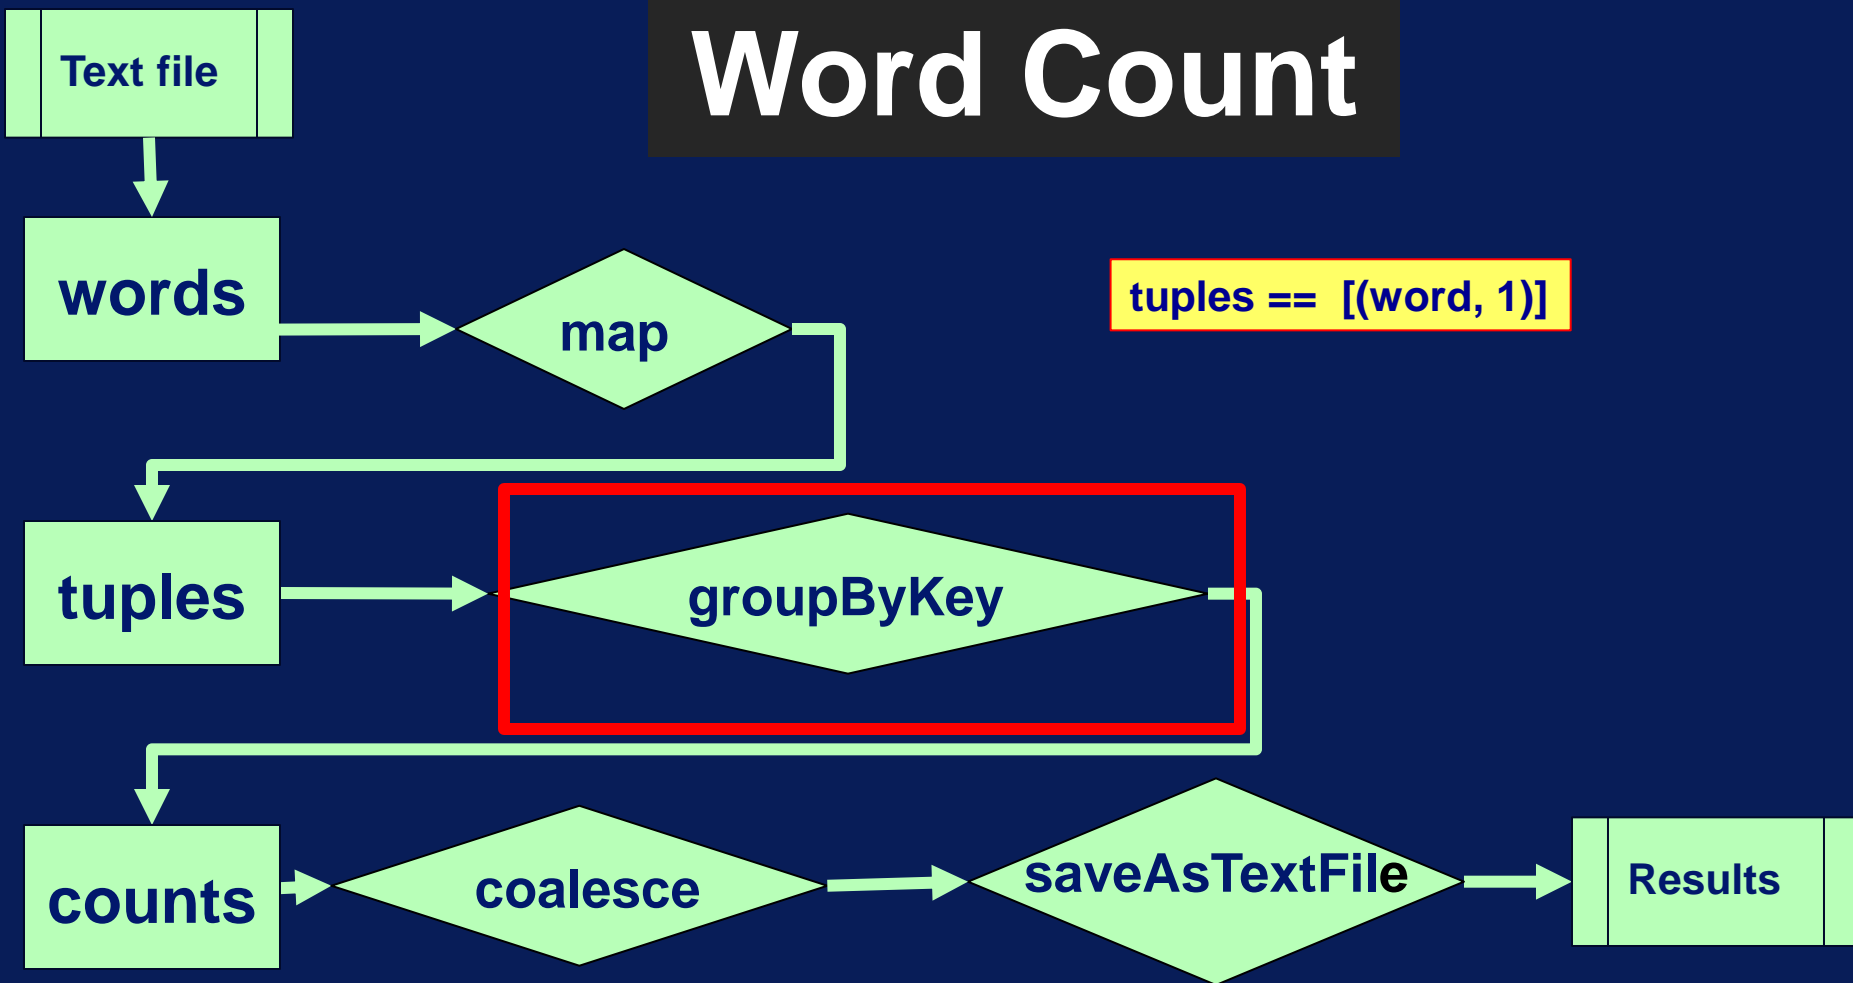


Wide Transformations

Word Count

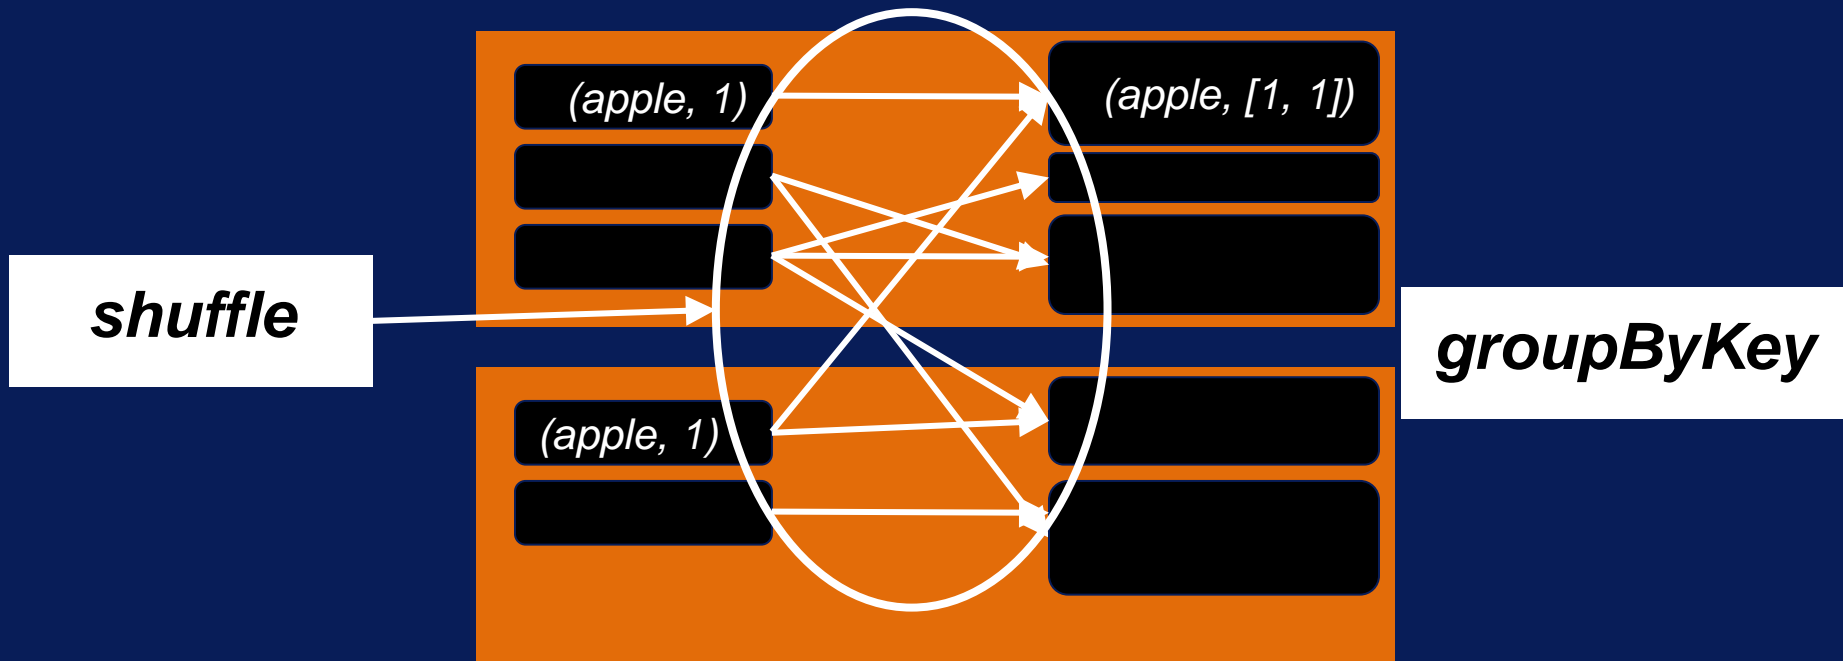


Word Count

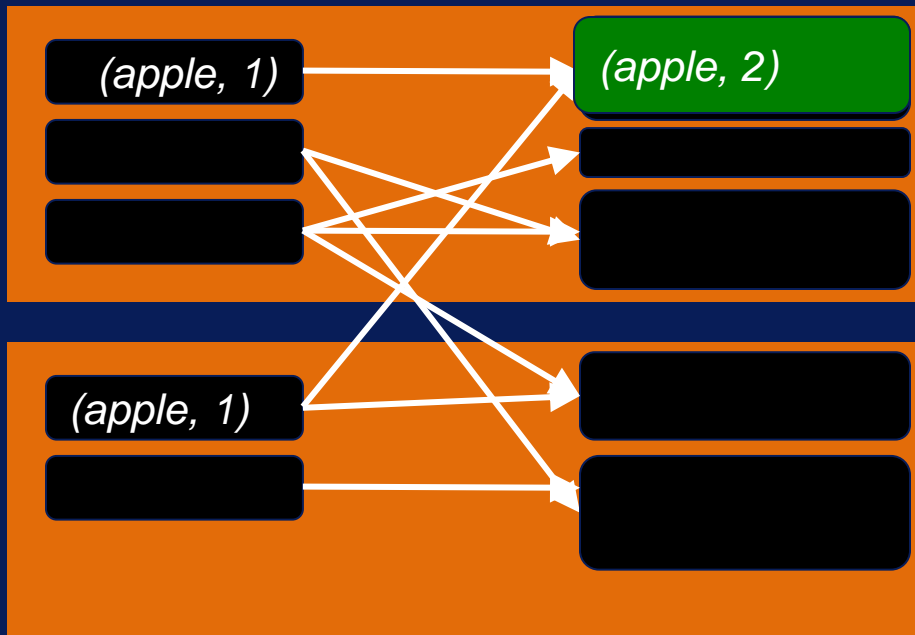


groupByKey


groupByKey : (K, V) pairs => (K, list of all V)



groupByKey + reduce



reduceByKey



Narrow

vs

Wide

map

groupByKey

(apple, 1)

(apple, [1, 1])

(apple, 1)

Many more transformations...

Full list of transformations at:

<https://spark.apache.org/docs/1.2.0/programming-guide.html#transformations>

Spark Core: Actions



After this video you will be able to..

- Explain the steps of a Spark pipeline ending with a collect action
- List four common action operations in Spark

Driver Program

Spark
Context



Spark
Context

flatMap


map

groupByKey

collect



Some Common Actions


Action	Usage
collect()	Copy all elements to the driver
take(n)	Copy first n elements
reduce(func)	Aggregate elements with func (takes 2 elements, returns 1)
	Save to local file or HDFS

Spark SQL



After this video you will be able to..

- Process structured data using Spark's SQL module
- Explain the numerous benefits of Spark SQL



Spark SQL

- Enables querying structured and unstructured data through Spark
- Provides a common query language
- Has APIs for Scala, Java and Python to convert results into RDDs


Relational Operations

Perform Relational Processing
such as Declarative Queries

Embed SQL queries
inside Spark
Programs

Business Intelligence Tools

Spark SQL connects to all BI tools that support JDBC or ODBC standard



<http://spark.apache.org/>

DataFrames

Distributed Data organized as
named columns

Look just like a
table in relational
databases

How to go Relational in Spark ?

Step 1: Create a SQLContext

```
from pyspark.sql import SQLContext  
sqlContext = SQLContext(sc)
```

How to go Relational in Spark ?

Create a DataFrame from

- an existing RDD
- a Hive table
- data sources

JSON → DataFrame

Read

```
df = sqlContext.read.json("/filename.json")
```

Display

```
df.show()
```

RDD of Row objects → DataFrame

Read

```
from pyspark.sql import SQLContext, Row
sqlContext = SQLContext(sc)
```

Load a text file and convert each line to a Row.

```
lines = sc.textFile("filename.txt")
cols = lines.map(lambda l: l.split(","))
data = cols.map(lambda p: Row(name=p[0], zip=int(p[1])))
```

Create DataFrame

```
df = sqlContext.createDataFrame(data)
```

Register the DataFrame as a table

```
df.registerTempTable("table")
```

Run SQL

```
Output = sqlContext.sql("SELECT * FROM table WHERE ...")
```

DataFrames are just like tables

<http://spark.apache.org/>

Show the content of the DataFrame

df.**show**()

Print the schema

df.**printSchema**()

Select only the "X" column

df.**select**("X").show()

Select everybody, but increment the discount by 5%

df.**select**(df["name"], df["discount"] + 5).show()

Select people height greater than 4.0 ft

df.**filter**(df["height"] > 4.0).show()

Count people by zip

df.**groupBy**("zip").count().show()

Spark SQL

Relational on Spark


Connect to variety of databases

Deploy business intelligence tools over Spark

Spark MLlib

After this video you will be able to..

- Describe what MLlib is
- List main categories of techniques available in MLlib.
- Explain code segments containing MLlib algorithms.



Spark MLlib

- Scalable machine learning library
- Provides distributed implementations of common machine learning algorithms and utilities
- Has APIs for Scala, Java, Python, and R

MLlib Algorithms & Techniques

- Machine Learning
 - Classification, regression, clustering, etc.
 - Evaluation metrics
- Statistics
 - Summary statistics, sampling, etc.
- Utilities
 - Dimensionality reduction, transformation, etc.

MILib Example – Summary Statistics

- Compute column summary statistics

```
from pyspark.mllib.stat import Statistics
```

1

```
# Data as RDD of Vectors
```

```
dataMatrix = sc.parallelize([ [1, 2, 3], [4, 5, 6], [7, 8, 9], [10, 11, 12] ])
```

2

```
# Compute column summary statistics.
```

```
summary = Statistics.colStats(dataMatrix)
```

3

```
print(summary.mean())
```

```
print(summary.variance())
```

4

```
print(summary.numNonzeros())
```

MLlib Example – Classification

- Build decision tree model for classification

```
from pyspark.mllib.tree import DecisionTree, DecisionTreeModel  
from pyspark.mllib.util import MLUtils
```

Read and parse data

```
data = sc.textFile("data.txt")
```

Decision tree for classification

```
model = DecisionTree.trainClassifier  
      (parsedData, numClasses=2)  
print(model.toDebugString())  
model.save(sc, "decisionTreeModel")
```

1

2

3

4

5

6

MLlib Example – Clustering

- Build k-means model for clustering

```
from pyspark.mllib.clustering import KMeans, KMeansModel  
from numpy import array
```

2

Read and parse data

```
data = sc.textFile("data.txt")  
parsedData = data.map(lambda line:  
    array([float(x) for x in line.split(' ')]))
```

3

k-means model for clustering

```
clusters = Kmeans.train (parsedData, k=3)
```

4

```
print(clusters.centers)
```

5

1

Main Take-Aways


- MLlib is Spark's machine learning library.
 - Distributed implementations
- Main categories of algorithms and techniques:
 - Machine learning
 - Statistics
 - Utility for ML pipeline

Spark Streaming



After this video you will be able to..

- Summarize how Spark reads streaming data
- List several sources of streaming data supported by Spark
- Describe Spark's sliding windows



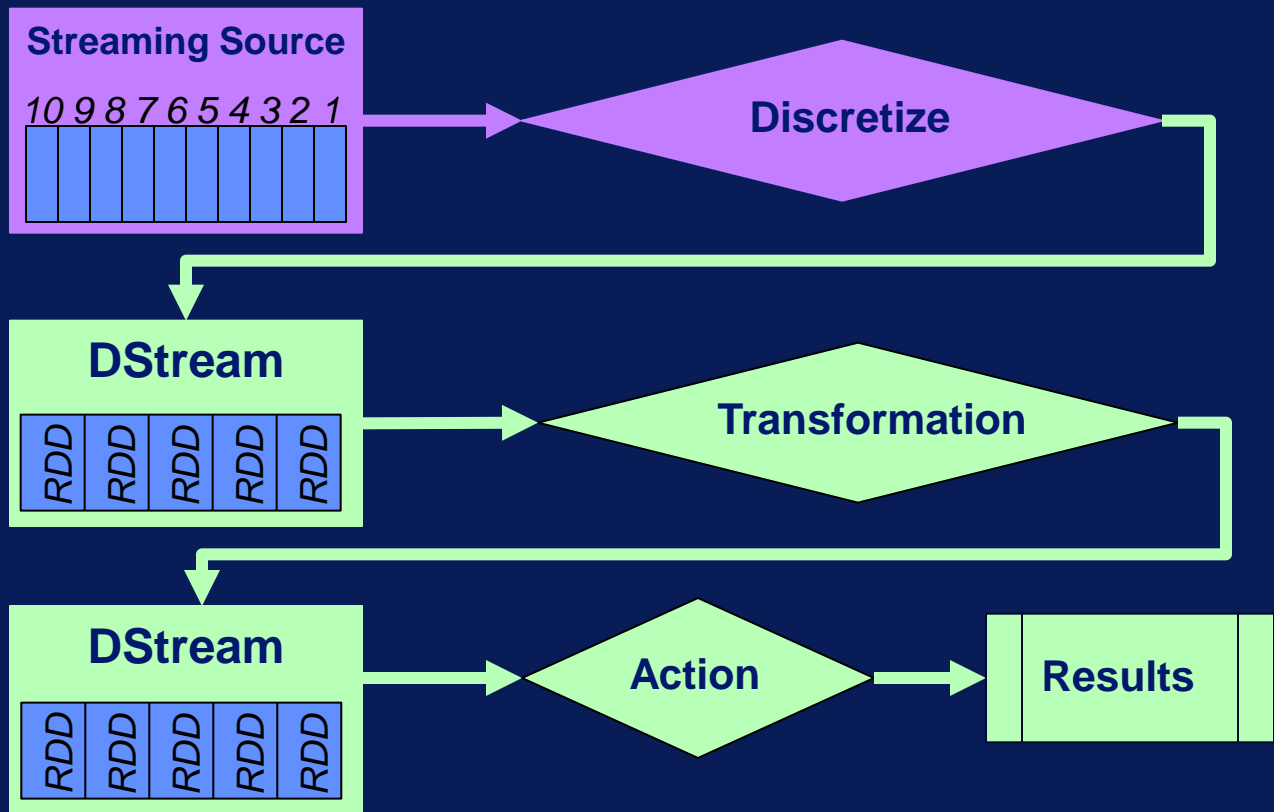
Spark Streaming

- Scalable processing for real-time analytics
- Data streams converted to discrete RDDs
- Has APIs for Scala, Java, and Python


Spark Streaming Sources

- Kafka
- Flume
- HDFS
- S3
- Twitter
- Socket
- ...etc.

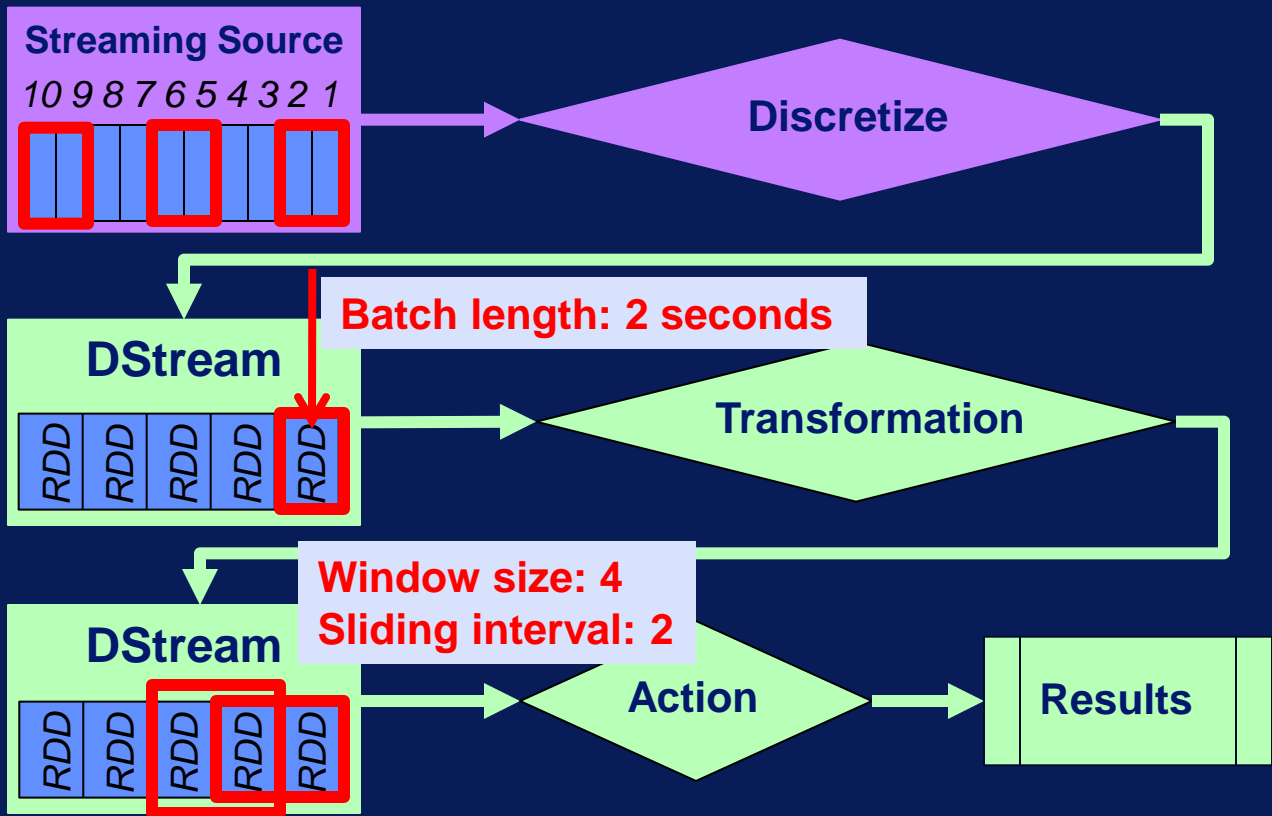
Creating and Processing DStreams



Creating and Processing DStreams




Creating and Processing DStreams



Main Take-Aways


- Spark uses DStreams to make discrete RDDs from streaming data.
 - Same transformations and calculations applied to batch RDDs can be applied
- DStreams can create a sliding window to perform calculations on a window of time.

Spark GraphX



After this video you will be able to..

- Describe what GraphX is
- Explain how Vertices and Edges are stored
- Describe how Pregel works at a high level



Spark GraphX

GraphX is Apache Spark's API for graphs and graph-parallel computation.

GraphX uses a property graph model.

**Both Nodes and
Edges can have
attributes and
values**

Properties → Tables

Vertex Table

Node properties

Edge Table

Edge properties


GraphX uses special RDDs

VertexRDD[A] extends RDD[(VertexID, A)]

EdgeRDD[ED, VD] extends RDD[Edge[ED]]

Triplets

The triplet view logically joins the vertex and edge properties.



Pregel API

Bulk-synchronous parallel messaging mechanism

Constrained to the topology of the graph

GraphX

Graph Parallel Computations

Special RDDs for storing Vertex and Edge information

Pregel operator works in a series of super steps